

External Undermining

Methodological memorandum — purpose, design and accountability of the dashboard

This memorandum is a standalone, self-explanatory accountability document accompanying the interactive “External Undermining” dashboard. It describes why the dashboard was built, the definition and methodology underlying it, the design choices made during its construction and the reasoning behind them, and the limitations that explicitly apply. It is intended to make every choice in the dashboard’s design — from the definition to the colour of a hexagon — traceable and justified, including for readers who did not see the dashboard take shape.

1. Purpose

The “External Undermining” dashboard is the third instrument in a series that makes the quality and resilience of EU democratic decision-making legible: alongside the EU Governance Assessment (EUGA, internal institutional erosion) and the Lobby Dashboard (legitimate but under-transparent influence from within). This third instrument documents a qualitatively different threat: attempts by non-EU actors to destabilise the European democratic process, influence public opinion, and undermine trust in politics.

The dashboard’s purpose is threefold: (1) to bring together a phenomenon that ordinarily plays out in the black box of intelligence services, national security assessments and scattered investigative journalism into a single accessible, searchable overview; (2) to structure that overview methodologically so that it never claims more than the underlying sources support; and (3) to thereby contribute to the recognition capacity of a citizen audience — who the actors are, which methods recur, which targets are typical — without itself reaching a judgement the sources do not warrant.

2. Why this is essential to the EU's core values

Article 2 TEU names democracy, the rule of law and respect for human rights as the values on which the Union is founded. Democracy in that sense presupposes that citizens can make their political choices freely and on an informed basis, without covert, illegitimate influence by actors standing outside their own democratic community. External undermining directly attacks that capacity: disinformation campaigns, covert financing of parties and campaigns, and cyber operations with a political aim are all directly intended to pollute the voter’s information environment and erode trust in democratic institutions.

For a broad public, this threat is largely invisible. National intelligence services publish annual reviews in their own language, the EEAS publishes annual FIMI reports at EU level, and investigative journalism documents individual cases — but these sources are rarely brought together in a way a citizen can actually search. The dashboard therefore serves a transparency function that is instrumental to democratic accountability: a citizen who can recognise the pattern of external influence — which actors, which networks, which targets recur — is better equipped to recognise new attempts when they occur. In that sense the dashboard is not a neutral reference work, but an instrument that contributes directly to the resilience that Article 2 TEU presupposes.

3. Definition and scope

For this dashboard, external undermining is defined as: activities by non-EU actors that, financially or otherwise, support parties or campaigns with a view to destabilising the democratic process within the EU, influencing public opinion, and undermining trust in politics.

3.1 What falls within scope

- Disinformation campaigns aimed at elections, referendums or public opinion generally.
- Covert or non-transparent financing of parties, campaigns or individual politicians by foreign actors.
- Network- and infrastructure-based operations (for example imitation media sites, bot networks) with a demonstrable political aim.
- Cyber operations with a politically electoral or society-disrupting component.
- Direct narrative interventions by state actors themselves, even without technical infrastructure (for example a public statement by a foreign intelligence service about a specific party).

3.2 What is deliberately excluded, and why

- **Purely domestic disinformation.** Even where harmful to public debate, no external actor is involved. Examples deliberately excluded from the count for this reason include the fabricated Tisza election-platform episode in Hungary (produced by the governing party itself), statements by Italian politicians echoing pro-Kremlin positions, and the Ibiza affair in Austria (a journalistic sting operation, not a genuine Russian actor).
- **Intra-EU interference.** The definition explicitly refers to non-EU actors. An allegation that an EU member state (Poland) interfered in Hungary's election campaign falls outside this definition and has therefore not been coded as an incident, even though such conduct — if proven — would be equally problematic. This is a deliberate boundary, not a substantive judgement about the seriousness of intra-EU interference.
- **International organisations without a collective defence function.** Organisations that happen to be headquartered on EU territory (for example the OPCW in The Hague) are excluded from having their own level in the dashboard. NATO is the exception: because the alliance collectively defends member states' territory, an attack on NATO is functionally connected to the security of EU member states themselves in a way that does not apply to other international organisations (see §6.2).
- **Physical/kinetic hybrid threats without a demonstrable political link.** GPS jamming and comparable physical incidents are, where they occur, attributed to the member state concerned rather than to NATO or the EU as a whole, unless a source itself draws a direct link to political influence.

4. Methodological framework

4.1 Dossier rather than index

For the EUGA series, a composite score was chosen, based on annual expert indices with full country coverage (V-Dem, World Justice Project). For external undermining, no comparable, fully covering expert index exists: the threat is covert, varies widely in form, and the available sources (intelligence services, the EEAS, investigative journalism) do not cover all member states equally well. A composite score would in that case suggest a precision the underlying data cannot support. A dossier approach was therefore chosen instead, modelled on the Lobby Dashboard: individual, source-traceable incidents rather than a single figure.

4.2 The analytical framework: four questions per dossier

Each country dossier consistently answers four questions, each with an explicit source basis, to prevent the dashboard itself from adding interpretation not found in the source:

- Scale — how many documented incidents, over what period.
- Targets — which election, institution, party or population group.
- Actors — which state or network the activity is attributed to, and with what degree of certainty.

- Aims — what the source itself describes as the intended effect, without interpretation beyond the source.

4.3 Two independent scales: severity and certainty

A recurring analytical insight during construction is that “striking” and “well-substantiated” are two different things that must not be conflated. The Slovakia dossier illustrates this sharply: the two viral deepfakes around the 2023 election carry a high severity rating but a low certainty rating (the creators were never identified), while the most certain case — a public statement by Russia’s own SVR — received a medium severity rating. Severity and certainty are therefore coded as two independent variables that are never merged.

4.4 Source hierarchy for the certainty rating

The certainty rating reflects the strength of the attribution, not the seriousness of the incident:

- **High** — attribution by a national security service, an EU institution (EEAS), or confirmed via multiple independent sources.
- **Medium** — attribution by a specialised research institute, cybersecurity firm or academic publication.
- **Low** — attribution by a single journalistic source or an unsubstantiated claim.

5. Key scope and design decisions

5.1 The EU institutional level as a separate entry

Alongside the 27 member states, the dashboard includes a separate “EU institutions” level, for cases where a source itself frames the target as an EU institution or an EU-wide process (for example the European Parliament elections as a whole), rather than a specific member state. This prevents such cases from being artificially attributed to a single member state, or lost between member-state dossiers.

5.2 NATO, but not other international organisations

At explicit request, it was reconsidered whether international organisations headquartered on EU territory deserve their own tile. The conclusion: NATO does, other organisations (such as the OPCW) do not. The reasoning: NATO collectively defends member states’ territory, so an attack on NATO bears directly on the security of EU member states themselves. An organisation that merely happens to be headquartered on EU territory, without that collective defence function, lacks that direct link. The NATO dossier has deliberately been kept the “thinnest” in the set: the available sources (EEAS reporting) describe the threat only in aggregate terms, without the incident-level detail available for member states, and that difference in source depth is stated explicitly within the dossier itself.

5.3 The symmetry principle: no separate bar for small countries or the Balkans

A recurring principle is that the same research effort and the same evidentiary bar apply to every country, regardless of size or region. This principle was tested in practice: for Croatia, initial searches yielded insufficient material, and — consistent with the principle — a weaker dossier was not accepted as a substitute; instead a second, targeted search round was carried out that did meet the standard. For Luxembourg, three search passes in two languages produced no qualifying incident; that outcome was not concealed or glossed over, but explicitly recorded as its own dossier type (see §5.4).

5.4 The distinction between “not yet examined” and “examined, none found”

The dashboard explicitly distinguishes between countries not yet examined and countries that have been examined but for which no qualifying incident was found. Luxembourg is the sole example of the latter category:

its dossier documents not an incident, but the search process itself (three search rounds, two languages, explicit comparison with neighbouring countries that are named in network reporting) and is included, with its own visual marker, as a clickable entry on the map. This prevents an absence of evidence from being mistaken for an absence of research.

5.5 Contested cases remain visibly contested

Where sources contradict one another, or where an attribution is itself disputed, the dashboard does not resolve the dispute in its own voice. Three examples: the annulled 2024 Romanian presidential election, where both the interference claim and the annulment itself are contested within Romania; the Hungary dossier, in which the reported Russian activity is said to support the incumbent government rather than the opposition — the reverse pattern from almost every other dossier — and which explicitly notes that a Polish role has also been alleged (and, per §3.2, not counted); and the Slovenia Black Cube dossier, in which the country's own intelligence service, SOVA, confirms foreign interference without confirming a state-level principal.

5.6 Automatic actor aggregation without manufactured precision

The actor summary on the dashboard's landing page is not maintained by hand, but is calculated automatically from the certainty-tagged data already present in each dossier. Incidents with a joint attribution (for example to Iran, China and Russia together) are placed in a dedicated "Multiple / joint actors" category rather than artificially split across individual countries. Incidents with no attributed actor (for example the Slovakia deepfakes) fall automatically into a "Not attributed" category. These residual categories are deliberately kept visible rather than engineered away.

6. Visualisation choices

6.1 Hexagonal grid cartogram

The map is a schematic grid cartogram with equal cell size per country, not a map with exact national borders. This is a deliberate trade-off between geographic recognisability and visual legibility: a grid cartogram prevents small countries (Malta, Luxembourg) from visually disappearing next to large ones (Germany, France), while the relative position of countries to one another remains recognisable. Switzerland is included as a recognisable, non-clickable "ghost hex" to signal that the gap in the grid pattern is a deliberate geographic gap, not a missing member state.

6.2 Colour coding: incident count, not severity

The colour intensity of each tile reflects the number of documented incidents for that country, not the average severity or certainty of those incidents. This is deliberate: a country with many low-severity incidents would otherwise appear visually heavier than a country with a single very high-severity incident, producing a distorted picture. Severity and certainty information is available per incident within the dossier view itself, not compressed into the map colour.

6.3 EU and NATO kept off the geographic grid

EU institutions and NATO are not geographic entities and were therefore, after an earlier version in which they sat among the member states, moved into their own, explicitly labelled block separate from the map. This change followed directly from user feedback that placement among member-state tiles created the impression of comparable, geographically situated units.

6.4 Two formats: the full dashboard and a compact teaser

Alongside the full, analytical dashboard, a compact teaser version exists, intended for inclusion in the accompanying article, with a click-through to the full dashboard. The teaser shows the same map and actor summary in condensed form, within the footprint of half an A4 page, and is updated quietly (without a new design round) whenever the dataset is extended, unless a substantive or design issue arises that requires a separate decision.

7. Limitations and caveats

The following limitations apply to the dashboard as a whole and are stated here explicitly, in addition to the dossier-specific caveats already included within the individual dossiers.

- **Publication and detection bias.** Countries with an active national reporting practice (Germany, France, Estonia, Sweden) show systematically more incidents than countries without a comparable practice. This reflects the capacity and willingness of sources to publish, not necessarily the true scale of the threat per country.
- **Russia's dominance in the aggregate figures.** Russia is the attributed actor in roughly four out of five coded incidents. This may be a genuine reflection of the threat, but could equally result from which actor is most actively tracked by Western intelligence services and investigative journalism. The dashboard makes no claim either way.
- **A snapshot in time.** The dashboard never claims completeness. New sources, further research or revised assessments may cause existing dossiers to grow or change in colour.
- **Contested attributions remain contested.** Where sources contradict one another, the dashboard does not choose between the readings on the authors' own authority; both are presented as the sources themselves present them.
- **Thin dossiers are not evidence of low risk.** Dossiers with few incidents (Malta, Cyprus, Luxembourg) reflect limited public documentation at the time of coding, not necessarily limited actual risk.
- **A schematic map, not a legal or geographic claim.** The grid cartogram makes no claim about borders, territory or sovereignty; cell positions are chosen solely for visual recognisability.

8. Maintenance and updating

The recommended update cycle is annual, following publication of the latest EEAS annual report on FIMI threats, supplemented by ongoing monitoring of national sources. New incidents are added to the underlying Excel incident database; the dashboard view and the teaser are then updated quietly (counts, map colours, actor summary) without requiring a new design or verification round with the commissioning party, unless a substantive or visual issue arises that calls for a separate decision.

9. Sourcing

Every incident in the dashboard is directly traceable to a source with name, publication date and URL, recorded in the accompanying Excel incident database ('Incidents' tab). Sources include national intelligence and security services, the European External Action Service (EEAS/EUvsDisinfo), specialised research institutes (including DFRLab, ISD, CSD, the NATO StratCom Centre of Excellence), national investigative journalism and, where applicable, judicial or parliamentary findings.

Aantekeningen in de marge — methodological memorandum for the External Undermining dashboard. Compiled by Aantekeningen in de marge, with extensive editorial and visual input from Claude, Anthropic's AI assistant.